

NIB 7072: Machine Learning Coursework Report

Student Name: Madushika Pramod Senevirathna

Student ID: COMScDS252P-002

[GitHub Repo](#)

Question 1: Associative Classification of Lower or Uncertain Self-Perceived Academic Success Using a 2024–2025 Global Higher-Education ChatGPT Survey

Abstract

This study uses the 2024–2025 wave of the global *Higher Education Students' Evolving Perceptions of ChatGPT* survey to evaluate whether other responses collected in the same survey can distinguish respondents who report lower or uncertain self-perceived academic success from those who report higher self-perceived success. The raw dataset contains 22,963 responses across 180 variables, collected in 7 languages from 120 countries and territories. A binary outcome target was derived from question Q35a ("I am successful in my studies"): responses 1–3 represent the positive class (lower or uncertain self-perceived success), while responses 4–5 represent the negative class (higher success). After cleaning duplicate entries and removing records with missing outcomes, a final labelled dataset of 17,426 records was established. The analysis does not estimate changes in perceptions over time, objective academic performance, or causal effects of ChatGPT use.

A robust preprocessing pipeline was implemented to prevent data leakage. This included four domain-specific composite features, ANOVA percentile feature selection, five-fold stratified cross-validation, and randomized hyperparameter search. Five distinct machine learning model families were evaluated: Logistic Regression, K-Nearest Neighbours (KNN), Random Forest, Histogram Gradient Boosting, and a Multilayer Perceptron (Artificial Neural Network), selecting positive-class F1-score as the primary metric. Logistic Regression achieved the highest cross-validated F1-score (0.6011 ± 0.0107) during training and was chosen as the winning model prior to testing. On the locked test set, Logistic Regression achieved an overall accuracy of 0.6546, precision of 0.5277, recall of 0.6429, F1-score of 0.5796, ROC-AUC of 0.7124, and PR-AUC of 0.5862. Although Histogram Gradient Boosting achieved higher overall accuracy and ranking scores, it yielded significantly lower recall for the positive class.

Permutation importance analysis showed that job confidence after graduation, first-year Bachelor's enrolment status, and emotion balance were the most influential predictor variables. These relationships represent empirical associations rather than causal effects. Subgroup fairness auditing showed minimal performance gaps across supported gender groups, but revealed substantial performance disparities across age brackets and survey languages.

Combined with the subjective nature of the survey outcome and convenience sampling, these limitations demonstrate that the model is unsuitable for automated academic decision-making.

Finally, the model was packaged and deployed as a production-grade REST API. An initial IAM-authenticated FastAPI application was built into a Docker container and deployed on Google Cloud Run. The service was updated to support public HTTPS access with built-in security controls, including a 65,536-byte request body limit and a sliding-window rate limit of 60 prediction requests per 60 seconds per instance. The deployment pipeline passed a suite of 27 local unit tests and was deployed via GitHub Actions (Run 29691386855, Git commit `b610a81a322264edc3c0490f9bdb8a42d6eb62b6`). Public endpoints `/health` and `/model-info` were verified with HTTP 200 responses, confirming successful cloud deployment.

1. Introduction

Generative artificial intelligence tools like ChatGPT have become increasingly popular among higher education students. However, AI adoption rates and perceived academic value vary significantly across academic disciplines, study levels, and geographic regions. Educational institutions are often interested in identifying factors linked to student academic experiences. However, building predictive models on survey data requires careful handling: survey responses are self-reported, cross-sectional, and often incomplete. Consequently, a machine learning model can identify statistical associations, but it cannot prove whether AI tool usage directly improves or hinders student academic performance.

Using only the 2024–2025 cross-sectional survey wave, this study builds an associative classifier focused on a single survey statement: Q35a ("I am successful in my studies"). It tests whether eligible responses drawn from Q6–Q34 can distinguish lower or uncertain self-perceived success from higher self-perceived success. The positive target class represents ratings 1–3, not objective academic failure, a diagnosis, or a verified need for intervention. The analysis does not examine how students' perceptions of ChatGPT evolved over time. It addresses three core questions:

1. How accurately can eligible Q6–Q34 responses collected in the same survey distinguish the two Q35a self-perceived-success groups?
2. What performance trade-offs exist across linear, distance-based, tree-based, boosting, and neural network algorithms?
3. Which survey features influence model predictions, and how does error behaviour vary across the audited demographic subgroups?

The purpose is to benchmark five leakage-controlled supervised classifiers for distinguishing lower or uncertain from higher contemporaneous self-perceived academic success, using positive-class F1 for model selection, and then examine feature associations and subgroup error patterns. This model is explicitly designed as an analytical benchmark rather than an early-warning intervention tool. In practical settings, asking students about their academic confidence directly is more transparent than trying to infer it from a lengthy questionnaire. Any future operational tool would need a shorter instrument administered at a prospectively defined prediction time, a meaningful later outcome, and independent educational, fairness, and ethical validation.

1.1 Related research

Existing literature in educational machine learning generally falls into two categories: research on student AI adoption patterns and predictive modelling of academic performance using institutional learning management system (LMS) data.

Study	Data and method	Relevant finding	Limitation for the present context
Ravšelj et al. (2025)	Large global student survey; descriptive analysis and ordinal regression	Perceptions and uses of ChatGPT vary across tasks and student contexts	Convenience sampling and self-report do not support representative or causal claims
Yağcı (2022)	1,854 students; six classical classifiers using midterm grade, department and faculty	Several model families predicted final-exam categories with useful accuracy	One Turkish-language course and a strong prior-grade predictor limit transferability
Li et al. (2022)	Multi-source campus behaviour; LSTM and two-dimensional CNN	Heterogeneous behavioural sources improved achievement classification	Institution-specific digital traces and objective outcomes differ from a global perception survey
Guanin-Fajardo, Guaña-Moya and Casillas (2024)	6,690 records from one Ecuadorian higher-education institution; academic and socioeconomic variables, feature selection, imbalance treatments and nine classifiers	XGBoost achieved accuracy 0.7949, F1 0.8306 and AUC 0.8775; a decision tree provided ten simpler rules	Evidence from one institution limits transferability, and the available variables did not support a full equity assessment
Rico-Juan, Cachero and Macià (2024)	322 Spanish Computer Engineering students across two mathematics subjects; prior performance, personality and Moodle engagement were modelled with several regressors, cross-validation and explanation methods	Prior performance improved on the baseline by 21%; adding personality and engagement increased the improvement to 27%, and risk groups were identified one month into the semester	The study used a small, subject-specific LMS dataset; reuse assumes a stable course context and some instructor machine-learning knowledge
Youssef et al. (2024)	353 UAE students; cross-sectional PLS-SEM	ChatGPT use was associated with engagement, critical	Small local sample and cross-sectional

Study	Data and method	Relevant finding	Limitation for the present context
		thinking and reported achievement	design cannot establish effects
Smerdon (2024)	Mixed-method study of permitted AI use in an undergraduate assessment	AI adopters did not achieve higher assignment scores; prior performance influenced adoption	One assessment and voluntary adoption constrain generalisation
Baek, Tate and Warschauer (2024)	1,001 US college students; regression and thematic analysis	Use varied with demographic, study and institutional-policy context	US self-report data do not establish transferability to multilingual settings

While previous studies have mostly focused on single-institution LMS data or regional surveys, this work systematically benchmarks five distinct machine learning model families on a large, multilingual dataset. Strict data leakage prevention is enforced by placing all feature selection and scaling inside validation folds, while combining model interpretability with fairness error auditing.

2. Data and Methods

2.1 Dataset and study design

Open-access survey data were obtained from *Higher Education Students’ Evolving Perceptions of ChatGPT: Global Survey Data from the Academic Year 2024–2025*, Mendeley Data Version 2 (Aristovnik et al., 2025; DOI [10.17632/nv2343nwsb.2](https://doi.org/10.17632/nv2343nwsb.2)). The dataset contains anonymous responses collected between October 2024 and February 2025 using convenience sampling. The survey was conducted in seven languages: Arabic, English, Hebrew, Italian, Japanese, Spanish, and Turkish. It covers student demographics, AI usage habits, perceived capabilities, ethical concerns, academic satisfaction, emotional states, and study performance.

The source title uses “evolving” because this survey repeated an earlier global survey. The present analysis uses only the 2024–2025 wave and neither compares survey waves nor tracks respondents over time.

The raw dataset comprised 22,963 rows and 180 columns. The mixture of Likert scales, binary choices, categorical demographics, and open text fields provides a realistic challenge for tabular machine learning workflows.

2.2 Outcome and predictor governance

The target variable was derived from question Q35a ("I am successful in my studies"), measured on a 5-point Likert scale. The binary target was defined as follows:

- **Class 1 (Positive Class — Lower or Uncertain Success):** Likert ratings 1, 2, or 3.
- **Class 0 (Negative Class — Higher Success):** Likert ratings 4 or 5.

Grouping rating 3 with ratings 1–2 operationalizes the pre-specified “lower or uncertain” class; it does not identify academic failure or a verified need for support. Records with missing Q35a values were removed rather than imputed.

Predictors were eligible responses drawn from Q6–Q34. Questions Q1–Q5 and Q35b–Q40 were excluded because of routing, privacy, or target-proxy leakage. Institution, free text, citizenship, country of study, gender, age, and survey language were also excluded from model features because of privacy, cardinality, or fairness concerns. Gender, age, and language were retained separately for fairness subgroup auditing. Individual items used in feature engineering were dropped after aggregation to eliminate feature redundancy.

2.3 Data quality and cleaning

The initial data audit identified several data quality challenges that required systematic preprocessing.

Finding	Observed value	Treatment
Raw dimensions	22,963 × 180	Preserved for the initial audit
Exact duplicate excess rows	127	Removed before splitting
Overall missing-cell rate	23.0413%	Handled by routing-aware exclusions and fold-fitted imputation
Rows with at least one missing value	22,403	Retained where the outcome was available
Raw non-missing Q35a responses before deduplication	17,427	17,426 labelled responses remained after duplicate removal
Institution values	8,662 unique	Excluded rather than high-dimensional encoding
Q13f free-text missingness	94.6348%	Free text excluded
Missing outcomes after deduplication	5,410	Removed; never imputed
Invalid coded questionnaire responses	0	No correction required

Minor data entry errors were corrected, such as converting text-formatted ages into numerical values and setting invalid age entries to missing. The age interquartile range (IQR) check flagged 1,490 values outside the 18–30 range; because these fell within valid survey bounds (18–100), they were preserved for demographic auditing. Age was not used as a model feature.

Missing values in the dataset were partly structural due to questionnaire skip logic. A missing indicator was added to allow models to recognize missing patterns. Removing 5,410 records due to missing targets reduced the dataset to 17,426 labelled respondents:

Outcome class	Count	Share
Higher self-perceived success (0)	10,971	62.958%
Lower/uncertain self-perceived success (1)	6,455	37.042%

The resulting class imbalance is moderate (37% positive class), making standard classification techniques suitable without synthetic oversampling. Exploratory data analysis confirmed key relationships across study fields, degree levels, and emotional indicators, as illustrated below.




Figure Q1.1. Missingness rates across top predictors and final binary outcome target distribution.




Figure Q1.2. Outcome distribution across degree levels, study fields, and learning modes.




Figure Q1.3. Spearman correlation matrix highlighting relationships between survey features and the target.

2.4 Composite feature construction

To reduce dimensionality and capture meaningful domain constructs, four continuous composite features were engineered:

Feature	Definition	Minimum coverage	Non-missing
ChatGPT task breadth rate	Proportion of answered Q18 task items used at least "Sometimes"	10 of 12 tasks	15,413
Capability mean	Mean of Q19a–Q19j perceived-capability items	8 of 10 items	15,424
Ethical-concern mean	Mean of Q22a–Q22j concern items	8 of 10 items	15,416
Emotion balance	Mean of positive-emotion items minus mean of negative-emotion items	6 of 7 items in each block	15,384

Task breadth ranges from 0 to 1, capability and ethical concern range from 1 to 5, and emotion balance ranges from –4 to +4 (where higher positive values reflect stronger positive affect). The

final pre-encoding dataset contained 98 input features: 10 nominal, 84 ordinal, and 4 engineered continuous features.

2.5 Train-test split and preprocessing

The dataset was split into an 80% training set and a 20% locked holdout test set using stratified random sampling with a fixed random seed.

Partition	Rows	Positive cases	Positive rate
Training	13,940	5,164	0.370445
Locked test	3,486	1,291	0.370338

To prevent data leakage, all learned preprocessing steps were encapsulated inside scikit-learn pipelines and fitted independently on each training fold during cross-validation:

- **Engineered Continuous Features:** Median imputation with missing value indicators, followed by Robust Scaling.
- **Ordinal Features:** Median imputation with missing value indicators, followed by Standard Scaling.
- **Nominal Features:** Missing value handling via a distinct category, followed by One-Hot Encoding.
- **Feature Selection:** Univariate ANOVA `SelectPercentile` (evaluating top 25%, 50%, and 75% feature sets).

For the winning Logistic Regression model, one-hot encoding expanded the feature space to 215 columns, of which ANOVA selection retained 107 features (50%). No synthetic oversampling (e.g., SMOTE) was used, as regularisation and class-weight balancing proved sufficient. The decision threshold was fixed at 0.50.

2.6 Model development and validation

Five complementary machine learning model families were benchmarked:

Model	Category	Rationale
Logistic Regression	Linear	Regularized, probabilistic and comparatively interpretable baseline
K-Nearest Neighbours	Distance based	Tests local similarity after scaling
Random Forest	Tree based	Represents nonlinearities and interactions with bagged trees
Histogram Gradient Boosting	Ensemble boosting	Efficient regularized boosting for medium-sized tabular data
Artificial Neural Network (multilayer perceptron)	Artificial neural network	Tests nonlinear representation learning with two hidden layers

Each algorithm was trained using 5-fold stratified cross-validation paired with a 6-candidate randomized hyperparameter search (30 fits per family, 150 fits total). Models were optimized using positive-class F1-score as the primary selection metric.

Model	Randomized search space	Selected deployment settings
Logistic Regression	ANOVA percentile: 25, 50, 75; C: 0.05, 0.2, 1.0, 5.0; penalty: L1, L2; class weight: none, balanced	ANOVA percentile 50; C 5.0; L1 penalty; balanced class weight
K-Nearest Neighbours	ANOVA percentile: 25, 50, 75; neighbours: 11, 21, 31, 51; weights: uniform, distance; distance power: 1, 2	Not selected for deployment
Random Forest	ANOVA percentile: 25, 50, 75; trees: 200, 350; maximum depth: none, 12, 20; minimum leaf size: 1, 3, 6; maximum features: square root, 0.5; class weight: none, balanced	Not selected for deployment
Histogram Gradient Boosting	ANOVA percentile: 25, 50, 75; learning rate: 0.03, 0.06, 0.10; iterations: 100, 180, 260; leaf nodes: 15, 31, 63; minimum leaf size: 15, 30, 60; L2 regularization: 0, 0.5, 2.0	Not selected for deployment
Artificial Neural Network	ANOVA percentile: 25, 50, 75; hidden layers: (64, 32), (96, 48), (128, 64); alpha: 0.0001, 0.001, 0.01; initial learning rate: 0.001, 0.003; batch size: 64, 128	Not selected for deployment

2.7 Explanation, error and fairness methods

Model interpretability was evaluated using Permutation Feature Importance computed on a representative sample of 1,200 test instances (3 random shuffles per feature). Partial Dependence Plots (PDP) and binned quantile summaries were also generated to analyze feature response curves.

Error analysis examined confusion matrix distributions, false positive vs. false negative rates, and high-confidence error cases (predicted probability ≥ 0.80). Fairness auditing evaluated True Positive Rate (TPR), False Positive Rate (FPR), and Selection Rate gaps across gender, age groups, and survey languages for subgroups with at least 100 observations.

3. Results

3.1 Cross-validation and model selection

The cross-validation results across all five candidate model families are summarized below:

Model	Accuracy	Precision	Recall	F1	F1 SD	ROC-AUC	PR-AUC
Logistic Regression	0.6691	0.5430	0.6731	0.6011	0.0107	0.7235	0.5936
Random Forest	0.6929	0.5856	0.5848	0.5850	0.0146	0.7373	0.6280
Histogram Gradient Boosting	0.7095	0.6556	0.4555	0.5370	0.0165	0.7437	0.6470
Artificial Neural Network (multilayer perceptron)	0.6926	0.6133	0.4715	0.5310	0.0201	0.7197	0.6132
K-Nearest Neighbours	0.6727	0.5692	0.4791	0.5201	0.0155	0.6906	0.5685

Logistic Regression achieved the highest positive-class F1-score (**0.6011**) and recall (**0.6731**). Histogram Gradient Boosting achieved the highest overall accuracy (0.7095) and ROC-AUC (0.7437), but suffered from lower recall (0.4555) at the default 0.50 threshold. Because the primary goal was capturing lower/uncertain success cases, Logistic Regression was selected as the winning model.

3.2 Locked-test comparison

All five trained models were evaluated on the locked 20% test set (3,486 records):

Model	Accuracy	Precision	Recall	F1	ROC-AUC	PR-AUC
Logistic Regression	0.6546	0.5277	0.6429	0.5796	0.7124	0.5862
Random Forest	0.6799	0.5673	0.5716	0.5694	0.7209	0.6230
Histogram Gradient Boosting	0.7022	0.6365	0.4570	0.5320	0.7344	0.6337
Artificial Neural Network (multilayer perceptron)	0.6925	0.6435	0.3803	0.4781	0.7116	0.6050
K-Nearest Neighbours	0.6810	0.5833	0.4857	0.5300	0.6925	0.5779




Figure Q1.4. Performance metrics across all five candidate algorithms on the locked holdout test set.

The test set results confirmed the cross-validation findings: Logistic Regression preserved the highest positive-class recall (0.6429) and F1-score (0.5796). For the winning model, higher-success classification achieved precision 0.7590, recall 0.6615, and F1 0.7069.

3.3 Error analysis

The confusion matrix for the winning Logistic Regression model on the test set is shown below:

	Predicted higher	Predicted lower/uncertain
Actual higher	1,452	743
Actual lower/uncertain	461	830




Figure Q1.5. Confusion matrix, ROC curve, and Precision-Recall curve for the selected Logistic Regression pipeline.

The model correctly classified 2,282 instances (65.46%). It generated 743 false positives (21.31%) and 461 false negatives (13.22%). Analyzing median feature profiles across prediction outcomes showed substantial overlap between correct and misclassified records:

Prediction outcome	Median age	Capability mean	Ethical-concern mean	Emotion balance	Confidence
Correct	21	3.7	3.0	1.000	0.674
False negative	22	3.7	3.1	1.143	0.608
False positive	21	3.5	3.0	0.571	0.617

Notably, 73 misclassified instances had predicted probabilities ≥ 0.80 , proving that high model confidence does not guarantee correct predictions.

3.4 Model explanation

Permutation feature importance identified the top features driving model predictions:

Rank	Input	Questionnaire meaning	Mean F1 decrease	SD
1	Q12	Confidence about obtaining a job after completing studies	0.0513	0.0015
2	Q9	First-year student in a Bachelor's degree	0.0189	0.0068
3	Emotion balance	Positive-emotion mean minus negative-emotion mean	0.0135	0.0081
4	Q10	Main field of study	0.0098	0.0027
5	Q25a	Perception that ChatGPT use is under the respondent's control	0.0091	0.0035




Figure Q1.6. Permutation feature importance (F1-score decrease) for top predictors.

Post-graduation job confidence (Q12) was the single strongest predictor. Emotion balance and first-year undergraduate status also contributed significantly to model predictions.

Partial dependence of influential engineered features




Figure Q1.7. Partial dependence plots for emotion balance and capability mean.

Binned analysis showed a clear negative trend between emotion balance and predicted positive probability:

Emotion-balance bin	N	Actual positive rate	Mean predicted probability
–4.001 to 0	756	0.4921	0.5814
0 to 0.571	483	0.4369	0.5386
0.571 to 1.143	637	0.3344	0.4744
1.143 to 1.857	624	0.3285	0.4236
1.857 to 4	594	0.2492	0.3553

Students with higher positive emotion balance reported significantly lower rates of academic uncertainty.

3.5 Subgroup audit

Fairness auditing revealed key demographic performance gaps across supported subgroups:

Audited attribute	Supported groups	TPR gap	FPR gap	Predicted-positive-rate gap
Gender	2	0.0408	0.0077	0.0026
Age band	4	0.2649	0.1912	0.2587
Survey language	5	0.2428	0.2109	0.1987

Gender performance was equal (selection rates of 0.4510 for female and 0.4484 for male). However, age gaps were severe: True Positive Rate dropped from 0.7093 for ages 18–21 down to 0.4444 for ages 30–39. Language gaps were similarly wide (TPR ranged from 0.7222 in Arabic to 0.4795 in Italian). These performance discrepancies reinforce that the model cannot be deployed for operational student profiling.

4. Deployment and MLOps Implementation


4.1 Validated model package

A production-ready deployment package was created under the `cloud-deployment` directory containing:

- `model.joblib` : Serialized scikit-learn pipeline (preprocessing, feature selection, and Logistic Regression).
- `feature_schema.json` : JSON schema defining input types, ranges, feature ordering, and composite rules.
- `metadata.json` : Model metadata, versioning (`q1-v1`), training metrics, and operational boundaries.
- `app.py` : FastAPI Web application implementing `/health` , `/model-info` , and `/predict` endpoints.
- `requirements.txt` , `Dockerfile` , and `api_contract.json` : Environment dependencies and container configuration.

Local testing verified that the API successfully reloads `model.joblib` , validates raw input JSON payloads, constructs composite features on the fly, and returns correct predictions matching the notebook pipeline. A suite of 27 automated unit tests passed locally.

4.2 Cloud deployment evidence and architecture



The microservice was deployed to **Google Cloud Run** using a fully automated GitHub Actions CI/CD workflow ([Run 29691386855](#), Git commit `b610a81a322264edc3c0490f9bdb8a42d6eb62b6`).

The deployed Cloud Run service (`q1-success-api`) is live at <https://q1-success-api-608121463228.asia-south1.run.app>. The container includes custom ASGI middleware that enforces two critical security controls:

1. Rejects incoming HTTP request payloads larger than **65,536 bytes (64 KiB)**.
2. Limits inference requests to **60 prediction requests per 60 seconds** per instance (returning HTTP 429 when exceeded).

Unauthenticated requests to public endpoints `/health` and `/model-info` returned HTTP 200 responses, confirming successful cloud operation.

4.3 CI/CD, monitoring and governance

The automated GitHub Actions workflow executes the following deployment steps:

1. Sets up the Python 3.12 environment and runs 27 unit tests.
2. Builds a lightweight Docker container image tagged with the Git commit hash.
3. Authenticates securely to Google Cloud using Workload Identity Federation (no hardcoded service account keys).
4. Pushes the container image to Artifact Registry and deploys to Cloud Run with scale-to-zero autoscaling (1 CPU, 1 GiB RAM, 0-3 instances).
5. Verifies service health via automated GET checks on `/health` and `/model-info` .

For real-world governance, future enhancements should incorporate Vertex AI model monitoring for feature drift, Cloud Armor WAF rate-limiting, and human-in-the-loop review processes.

5. Discussion

5.1 Interpretation and model suitability

The experiments demonstrate moderate predictive signal but confirm that survey-based machine learning is not ready for automated academic interventions. Logistic Regression proved to be an efficient baseline due to fast execution, small memory footprint, and clear coefficient interpretability.

However, a precision of 0.5277 means nearly half of the students flagged as "lower/uncertain success" actually reported higher success. Furthermore, a recall of 0.6429 misses over one-third of students who express academic uncertainty. These error rates, combined with 73 high-confidence misclassifications, rule out using this system for automated student grading, academic placement, or administrative tracking.

5.2 Limitations

1. **Single-Wave, Cross-Sectional Self-Reports:** Outcome Q35a and the eligible Q6–Q34 predictors originate from the same 2024–2025 survey snapshot. Despite the source dataset's title, the model does not measure evolving perceptions, compare survey waves, forecast a later objective outcome, or establish causal direction.
2. **Convenience Sampling Bias:** The survey relied on voluntary online sampling, limiting global statistical representativeness.
3. **Subgroup Disparities:** Significant accuracy gaps across age and language groups require prospective calibration before practical use.

6. Conclusion

A benchmark of five machine learning model families was completed on 17,426 student survey responses. Logistic Regression was selected via 5-fold cross-validation and achieved a test set F1-score of 0.5796, recall of 0.6429, and ROC-AUC of 0.7124. The model was containerized and successfully deployed to Google Cloud Run with custom rate-limiting and payload size security controls. While the pipeline serves as a robust methodological benchmark, high error rates and demographic performance gaps confirm that self-reported survey models must not be used for autonomous academic decision-making.

References

- Aristovnik, A., et al. (2025). *Higher Education Students' Evolving Perceptions of ChatGPT: Global Survey Data from the Academic Year 2024–2025* (Version 2) [Dataset]. Mendeley Data. <https://doi.org/10.17632/nv2343nwsb.2>
- Baek, C., Tate, T., & Warschauer, M. (2024). "ChatGPT seems too good to be true": College students' use and perceptions of generative AI. *Computers and Education: Artificial Intelligence*, 7, 100294. <https://doi.org/10.1016/j.caeai.2024.100294>
- Guanin-Fajardo, J. H., Guaña-Moya, J., & Casillas, J. (2024). Predicting academic success of college students using machine learning techniques. *Data*, 9(4), 60. <https://doi.org/10.3390/data9040060>
- Li, X., Zhang, Y., Cheng, H., Li, M., & Yin, B. (2022). Student achievement prediction using deep neural network from multi-source campus data. *Complex & Intelligent Systems*, 8, 5143–5156. <https://doi.org/10.1007/s40747-022-00731-8>
- Ravšelj, D., Keržič, D., Tomaževič, N., Umek, L., Brezovar, N., et al. (2025). Higher education students' perceptions of ChatGPT: A global study of early reactions. *PLOS ONE*, 20(2), e0315011. <https://doi.org/10.1371/journal.pone.0315011>

Rico-Juan, J. R., Cachero, C., & Macià, H. (2024). Study regarding the influence of a student's personality and an LMS usage profile on learning performance using machine learning techniques. *Applied Intelligence*, 54, 6175–6197. <https://doi.org/10.1007/s10489-024-05483-1>

Smerdon, D. (2024). AI in essay-based assessment: Student adoption, usage, and performance. *Computers and Education: Artificial Intelligence*, 7, 100288. <https://doi.org/10.1016/j.caeai.2024.100288>

Yağcı, M. (2022). Educational data mining: Prediction of students' academic performance using machine learning algorithms. *Smart Learning Environments*, 9, 11. <https://doi.org/10.1186/s40561-022-00192-z>

Youssef, E., Medhat, M., Abdellatif, S., & Al Malek, M. (2024). Examining the effect of ChatGPT usage on students' academic learning and achievement: A survey-based study in Ajman, UAE. *Computers and Education: Artificial Intelligence*, 7, 100316. <https://doi.org/10.1016/j.caeai.2024.100316>

Question 2: Brent Crude Oil Price Forecasting and Volatility Analysis

Abstract

This study examines Brent crude oil price forecasting and return volatility. It uses 5,315 observed prices from 4 January 2005 to 31 December 2025. ARIMA, XGBoost and LSTM forecast the price 30 observed price days ahead and are compared with naive persistence. A Student-t GARCH(1,1) model separately forecasts return volatility one observed-price interval ahead.

The modelling process keeps the data in time order and uses target-date gaps at split boundaries. All price models are tested on the same 784 held-out forecast origins. ARIMA(2,0,0) has the lowest test RMSE at 6.419 USD per barrel, but its improvement over persistence is only 2.57%. XGBoost and LSTM do not improve on the naive benchmark.

The return series shows strong conditional heteroscedasticity. GARCH estimates highly persistent volatility, with $\alpha + \beta = 0.9922$, although its correlations with observed volatility proxies are weak. The models are therefore treated as transparent planning and risk benchmarks, not as autonomous trading systems.

1. Study Design and Data

1.1 Data source and scope

The analysis uses the Federal Reserve Bank of St. Louis series **DCOILBRENTU**, Europe Brent Spot Price FOB, whose original source is the U.S. Energy Information Administration (EIA). The series is quoted in U.S. dollars per barrel, published daily, and not seasonally adjusted. Brent is appropriate for this analysis because the sample contains long trends, changing market

regimes, abrupt shocks, and volatility clustering. These properties permit both price forecasting and a separate examination of conditional return variance.

The fixed data window is 1 January 2005 to 31 December 2025. The first available price is 4 January 2005. The downloaded data are retained only after schema and date-coverage checks; the filtered data used in the executed notebook have SHA-256 hash

5b79068704c8b14cf8482ad3ab8dc5a53dee022731d408569960ebeb54cdda6 .

Property	Executed value
Observed date range	2005-01-04 to 2025-12-31
Business-weekday positions	5,478
Observed prices	5,315
Unavailable weekday prices	163
Duplicate dates after cleaning	0
Mean price	75.692 USD/bbl
Standard deviation	23.656 USD/bbl
Minimum	9.120 USD/bbl
First quartile	58.555 USD/bbl
Median	72.560 USD/bbl
Third quartile	91.980 USD/bbl
Maximum	143.950 USD/bbl

1.2 Forecasting targets

The price models estimate the direct 30-observed-price-day target

$$y_t = P_{t+30},$$

where P_t is the Brent spot price observed at forecast origin t . An “observed price day” means the next row with a published FRED price; it is not an interpolated calendar or business day.

For volatility analysis, consecutive observed prices are converted to percentage log returns:

$$r_t = 100 [\log(P_t) - \log(P_{t-1})] .$$

GARCH estimates the conditional variance and standard deviation of r_t one observed-price interval ahead. It is not included in the price-model ranking because volatility and price level are different targets.

1.3 Literature context and model rationale

Recent research demonstrates that crude oil price forecasting heavily depends on model architecture, predictor selection, validation design, and market regime stability.

Study	Main focus	Relevance and limitation for this study
Ziółkowski (2024)	Compared LSTM, Prophet and XGBoost for WTI and Brent	Supports a common out-of-sample comparison of different model families.
Jin and Xu (2024)	Applied neural networks to several energy prices	Shows the flexibility of neural forecasting, but careful configuration and validation are still required.
Jha et al. (2024)	Used classical machine learning with multivariate crude-oil information	Shows the value of external predictors, which the present univariate design cannot reproduce.
Alruqimi and Di Persio (2024)	Studied multi-step Brent forecasting with recurrent networks	Highlights the difficulty of multi-step forecasting and the sensitivity of deep models to design choices.
Zhang et al. (2023)	Compared OVX, GARCH and stochastic-volatility forecasts for Brent and WTI crude-oil futures	Provides peer-reviewed evidence that volatility forecasts should be compared against observable market information and alternative variance models; the present study is narrower because it estimates only a univariate GARCH benchmark.
Chung (2024)	Compared GARCH-family and machine-learning volatility methods in a preprint	Shows that energy-market variance can be persistent but remains difficult to forecast accurately.

Based on these findings, a complementary model set was selected: ARIMA serves as an interpretable linear time-series baseline, XGBoost captures non-linear feature interactions, and LSTM models sequential dependencies. All models are compared against a naive persistence benchmark to test whether added computational complexity yields real-world predictive gains. In parallel, a GARCH model estimates conditional volatility rather than price direction.

2. Exploratory Analysis and Preprocessing

2.1 Missing observations and extreme returns

Data cleaning audited 163 missing business weekday prices. Instead of applying forward-fill or linear interpolation—which would artificially smooth return volatility—only published quote dates were preserved.

Across 5,314 calculated percentage log returns, modified z-score analysis ($|z| > 5$) identified 40 extreme return observations. The most severe single-day movements occurred during the April 2020 market crash (-64.37% on April 21, 2020, followed by +41.20% on April 22, 2020). Because these represent true market shocks rather than data errors, they were retained in the dataset. Robust feature scaling was applied for LSTM, while Student-t error distribution was specified for GARCH to handle heavy tails.

2.2 Trend, seasonality, and irregular variation

Plotting Brent crude oil prices and rolling averages shows strong non-stationary trend behavior and shifting volatility regimes.




Figure 1. Daily Brent crude oil prices with 30-day and 252-day moving averages (2005–2025).




Figure 2. Percentage log returns, highlighted extreme outliers, and rolling standard deviation.

Robust Seasonal-Trend Decomposition using LOESS (STL) was performed on monthly average prices with a 12-month seasonal cycle:




Figure 3. Robust STL decomposition of monthly average Brent prices.

STL component	Standard deviation
Trend	18.994
Seasonal	2.936
Residual	10.251

The decomposition confirms that long-term trend variation (SD 18.994) dominates the series, whereas seasonality is minor (SD 2.936).

2.3 Temporal features

A total of 24 temporal predictor features were engineered using information available at or before forecast origin t :

- Current price and current log return;
- Price lags at 1, 2, 5, 10, 20, 30, and 60 observations;
- Return lags at 1, 2, 5, 10, and 20 observations;
- Rolling price means (5, 20, 60 observations);
- Rolling return volatility (5, 20, 60 standard deviations);

- Cyclical sine/cosine transformations for month and day of year.

The prepared supervised dataset comprised 5,225 rows spanning forecast origins from March 30, 2005 to November 17, 2025.

2.4 Stationarity and autocorrelation

Augmented Dickey-Fuller (ADF) unit root tests were conducted to assess stationarity:

Series	ADF statistic	p-value	Lags	Decision at 5%
Price level	-2.5279	0.1088	6	Fail to reject unit root (Non-stationary)
First price difference	-31.3714	<0.0001	5	Reject unit root (Stationary)
Log return (%)	-11.9581	4.16×10^{-22}	32	Reject unit root (Stationary)

The raw price series is non-stationary, whereas differenced prices and log returns are stationary ($p < 0.0001$). Autocorrelation (ACF) and Partial Autocorrelation (PACF) plots of differenced prices show values close to zero past lag 1, indicating weak linear memory.




Figure 4. ACF and PACF plots for first-differenced Brent prices.

3. Price-Forecasting Methodology

3.1 Temporal validation design

To prevent data leakage in time-series forecasting, chronological purging was implemented with a 30-day target boundary gap between train, validation, and test splits:

Split	Rows	Forecast-origin range	Latest target date
Training	3,627	2005-03-30 to 2019-07-29	2019-09-09
Validation	754	2019-09-10 to 2022-08-26	2022-10-11
Test	784	2022-10-12 to 2025-11-17	2025-12-31

All models were evaluated on the same 784 test origins using Root Mean Squared Error (RMSE), Mean Absolute Error (MAE), Mean Absolute Percentage Error (MAPE), and Mean Absolute Scaled Error (MASE):

$$\text{MASE} = \frac{\frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|}{\frac{1}{T-1} \sum_{t=2}^T |P_t - P_{t-1}|}$$

The denominator represents the mean absolute 1-step price change on pre-test data (1.1363 USD/bbl). A MASE value greater than 1 indicates that the model's 30-step forecast error is larger than an average single-day price movement.

The relative percentage RMSE improvement over Naive Persistence is calculated as:

$$\text{RMSE reduction relative to persistence} = 100 \left(1 - \frac{\text{model RMSE}}{\text{persistence RMSE}} \right)$$

3.2 Naive benchmark

The persistence benchmark assumes that the price 30 days ahead equals the current price ($\hat{P}_{t+30} = P_t$). Any candidate machine learning model must outperform persistence to prove practical value.

3.3 ARIMA

A grid of 12 ARIMA candidates ($p \in \{0, 1, 2\}$, $d \in \{0, 1\}$, $q \in \{0, 1\}$) was evaluated across 30 rolling validation origins. **ARIMA(2,0,0)** achieved the lowest validation RMSE (10.776 USD/bbl) and was selected for test evaluation.

3.4 XGBoost

XGBoost was tuned using 12 hyperparameter configurations across 4 expanding `TimeSeriesSplit` folds with a 30-row purging gap. The winning parameters were:

- Number of trees: 150, Max depth: 2, Learning rate: 0.03, Subsample: 0.85, L1 reg: 0.10, L2 reg: 1.00.
- Cross-validation RMSE: 13.035 USD/bbl.

3.5 LSTM

The LSTM architecture comprised 1 LSTM layer (32 units, lookback 60), 10% dropout, a 16-unit ReLU dense layer, and a single output node. Model training used Adam optimization, MSE loss, batch size 32, and early stopping (15 epochs). Features were scaled using `RobustScaler` and targets using `MinMaxScaler` fitted strictly on training data.

3.6 Reproducibility

Random seeds were fixed at 42 across Python, NumPy, TensorFlow, and XGBoost. Full dependencies are documented in the companion notebook environment.

4. Price-Forecasting Results

4.1 Model selection evidence

Validation performance across tuning splits was recorded as follows:

Model and selected configuration	Validation MAE	Validation RMSE	Validation MAPE
ARIMA(2,0,0), 30 sampled origins	8.737	10.776	18.33%
XGBoost, full validation block	10.909	15.183	21.10%
LSTM, full validation block	11.798	15.043	23.70%

4.2 Held-out comparison

The holdout test set results (784 forecast origins) across all models are presented below:

Rank	Model	MAE (USD/bbl)	RMSE (USD/bbl)	MAPE	MASE	RMSE reduction relative to persistence (%)
1	ARIMA	5.223	6.419	6.73%	4.596	2.57%
2	Naive persistence	5.288	6.588	6.79%	4.654	0.00%
3	XGBoost	6.090	7.985	7.83%	5.359	-21.20%
4	LSTM	7.761	9.809	10.08%	6.830	-48.88%




Figure 5. Held-out 30-day price forecasts versus actual Brent spot prices.




Figure 6. Test set forecast error comparison (MAE, RMSE, MAPE).

ARIMA(2,0,0) achieved the lowest test RMSE (**6.419 USD/bbl**), outperforming Naïve Persistence by a modest **2.57%**. Complex non-linear models performed significantly worse: XGBoost increased RMSE by 21.20% over persistence, while LSTM increased RMSE by 48.88%. All MASE scores exceed 1, demonstrating the inherent difficulty of 30-day point forecasting in crude oil markets.

5. Conditional Volatility Modelling

5.1 Evidence of conditional heteroscedasticity

Financial return series frequently exhibit volatility clustering (periods of high volatility followed by high volatility). ARCH-LM test results strongly rejected the null hypothesis of constant variance:

Statistic	Value
ARCH-LM	970.433
ARCH-LM p-value	7.85×10^{-193}
F statistic	61.648
F-test p-value	1.69×10^{-218}




Figure 7. Rolling return volatility and squared return autocorrelation confirming ARCH effects.

5.2 GARCH(1,1)-Student-t estimates

A constant-mean GARCH(1,1) model was fitted with Student-t innovation errors:

$$\sigma_t^2 = \omega + \alpha \epsilon_{t-1}^2 + \beta \sigma_{t-1}^2$$

Parameter	Estimate	Interpretation
μ	0.0730	Conditional mean return (%)
ω	0.059428	Baseline variance component
α	0.085264	Immediate shock reaction coefficient
β	0.906976	Volatility persistence coefficient
$\alpha + \beta$	0.992239	Total volatility persistence
ν	6.101093	Student-t degrees of freedom (heavy tails)

The parameter sum $\alpha + \beta = 0.9922$ demonstrates high persistence in volatility shocks, while $\nu = 6.10$ confirms fat-tailed return distributions.

5.3 One-step forecast evaluation

Expanding 1-step volatility forecasts were evaluated against absolute and squared return proxies:

Comparison	Value	Unit
MAE: forecast volatility vs. absolute return	1.096	Percentage points
RMSE: forecast volatility vs. absolute return	1.331	Percentage points
MAE: forecast variance vs. squared return	4.012	Squared percentage points
Correlation: forecast volatility vs. absolute return	0.176	Unitless
Correlation: forecast variance vs. squared return	0.153	Unitless




Figure 8. Expanding 1-step GARCH volatility forecasts versus 20-day realized volatility proxy.

While GARCH successfully tracks broad volatility regime shifts, its correlation with noisy daily return proxies remains weak (0.176).

6. Discussion and Critical Reflection

The findings confirm that linear time-series methods (ARIMA) outperform complex non-linear machine learning models (XGBoost, LSTM) for 30-day crude oil price prediction. Tree ensembles struggles to extrapolate beyond historical price ranges, while LSTMs are prone to overshooting during sudden trend reversals.

However, ARIMA's 2.57% gain over Naive Persistence is marginal. Because crude oil markets are heavily influenced by exogenous geopolitical events, macroeconomic indicators, and OPEC decisions—which are omitted in univariate models—point forecasting past short horizons remains highly uncertain. GARCH volatility estimation provides a useful risk metric for portfolio management, but cannot be used as an independent automated trading signal.

7. Conclusion

Univariate forecasting models were evaluated on 5,315 daily Brent crude oil prices (2005–2025). ARIMA(2,0,0) achieved the lowest test RMSE (6.419 USD/bbl), outperforming Naive Persistence by 2.57%, whereas XGBoost and LSTM failed to outperform persistence. GARCH(1,1) confirmed high volatility persistence ($\alpha + \beta = 0.9922$). These results highlight that increased model complexity does not guarantee superior time-series forecasting, and univariate models should be treated as risk planning baselines rather than autonomous trading algorithms.

References

- Alruqimi, M., & Di Persio, L. (2024). Enhancing multi-step Brent oil price forecasting with ensemble multi-scenario Bi-GRU networks. *International Journal of Computational Intelligence Systems*, 17, Article 225. <https://doi.org/10.1007/s44196-024-00640-3>
- Chung, S. (2024). *Modelling and forecasting energy market volatility using GARCH and machine learning approach* [Preprint]. arXiv. <https://doi.org/10.48550/arXiv.2405.19849>
- Federal Reserve Bank of St. Louis. (2026). *Crude oil prices: Brent—Europe (DCOILBRETEU)* [Data set; original source: U.S. Energy Information Administration]. <https://fred.stlouisfed.org/series/DCOILBRETEU>
- Hyndman, R. J., & Koehler, A. B. (2006). Another look at measures of forecast accuracy. *International Journal of Forecasting*, 22(4), 679–688. <https://doi.org/10.1016/j.ijforecast.2006.03.001>
- Jha, N., Tanneru, H. K., Palla, S., & Mafat, I. H. (2024). Multivariate analysis and forecasting of the crude oil prices: Part I—Classical machine learning approaches. *Energy*, 296, 131185. <https://doi.org/10.1016/j.energy.2024.131185>
- Jin, B., & Xu, X. (2024). Price forecasting through neural networks for crude oil, heating oil, and natural gas. *Measurement: Energy*, 1, 100001. <https://doi.org/10.1016/j.meane.2024.100001>

U.S. Energy Information Administration. (2026). *Europe Brent Spot Price FOB (dollars per barrel)*. <https://www.eia.gov/dnav/pet/hist/rbrted.htm>

Zhang, Z., Raza, M. Y., Wang, W., & Sui, L. (2023). Volatility predictability in crude oil futures: Evidence based on OVX, GARCH and stochastic volatility models. *Energy Strategy Reviews*, 50, 101209. <https://doi.org/10.1016/j.esr.2023.101209>

Ziółkowski, K. (2024). Forecasting WTI & Brent crude oil price using LSTM, Prophet and XGBoost—Comparative analysis. In N. T. Nguyen et al. (Eds.), *Recent challenges in intelligent information and database systems* (CCIS Vol. 2145, pp. 171–181). Springer. https://doi.org/10.1007/978-981-97-5934-7_15

Question 3: Job-Shop Scheduling with a Genetic Algorithm and Mixed Integer Programming

1. Problem Definition and Literature Review

1.1 Prescriptive decision and objective

Job-shop scheduling is a prescriptive analytics problem: given jobs, machines, processing times, and operating constraints, the decision is **what processing sequence should be used**. This aligns with the framing of optimization as recommending actions under constraints, and of genetic algorithms (GAs) as metaheuristics for complex scheduling problems.

This analysis uses the Fisher–Thompson FT06 benchmark as a controlled proxy for a manufacturing job shop. Six jobs must each complete six ordered operations on six machines. An operation cannot be interrupted, every job must follow its specified route, and a machine can process only one operation at a time. The objective is to minimize the makespan, the completion time of the final operation:

$$\min C_{\max}.$$

Reducing makespan improves throughput and shortens the time required to finish the production batch. FT06 is not data from a particular factory; it is a standard benchmark that isolates the sequencing decision before real-world additions such as due dates, breakdowns, and setup times.

A practical example is a small manufacturing workshop with six customer orders and six machines. Each order must pass through six machining stages in a fixed order, while each machine can process only one operation at a time. The production planner must decide which available operation each machine should process next. Reducing the makespan allows the workshop to complete the full batch sooner. FT06 represents this decision in a controlled form, although a real factory would also need to consider due dates, setup times, maintenance, staff availability and machine breakdowns.

1.2 Mathematical formulation

The model is presented in the standard order of decision variables, objective, constraints, and domains. Let $\mathcal{O}$ be the set of operations, $\mathcal{P}$ the set of consecutive-operation pairs belonging to the same job, and $\mathcal{Q}$ the set of unordered pairs of operations requiring the same machine. For operation $a \in \mathcal{O}$, p_a is its processing time and μ_a its machine. The scheduling horizon is $H = 197$, the sum of all processing times.

Decision variables are:

- $S_a \geq 0$: start time of operation a ;
- $C_a \geq 0$: completion time of operation a ;
- $C_{\max} \geq 0$: makespan;
- $x_{ab} \in \{0, 1\}$ for $\{a, b\} \in \mathcal{Q}$: 1 when a is scheduled before b , and 0 otherwise.

The objective is:

$$\min C_{\max}.$$

Completion-time **equality constraints** are:

$$C_a = S_a + p_a, \quad a \in \mathcal{O}.$$

Job-precedence **inequality constraints** are:

$$C_a \leq S_b, \quad (a, b) \in \mathcal{P}.$$

For each same-machine pair $\{a, b\} \in \mathcal{Q}$, the binary variable activates one of two non-overlap inequalities:

$$C_a \leq S_b + H(1 - x_{ab}),$$

$$C_b \leq S_a + Hx_{ab}.$$

The makespan covers every operation:

$$C_a \leq C_{\max}, \quad a \in \mathcal{O}.$$

Together with $S_a, C_a, C_{\max} \geq 0$ and $x_{ab} \in \{0, 1\}$, these constraints define a feasible job-shop schedule. Machine assignments are fixed benchmark parameters, so separate assignment variables are unnecessary.

1.3 Literature review and method selection

Study	Contribution and relevance
Fisher and Thompson (1963)	Introduced influential job-shop test problems and probabilistic combinations of scheduling rules; FT06 originates from this benchmark family.
Garey, Johnson, and Sethi (1976)	Established the computational difficulty of job-shop makespan scheduling, motivating non-enumerative methods.
Adams, Balas, and Zawack (1988)	Developed the shifting-bottleneck heuristic, showing the value of problem-specific scheduling search.
Cheng, Gen, and Tsujimura (1996)	Reviewed GA representations for job-shop scheduling and emphasized that representation and operators must preserve meaningful schedules.
Gonçalves, Mendes, and Resende (2005)	Combined random-key GA search, schedule generation, and local search, demonstrating the strength of scheduling-specific hybrid GAs.
Ku and Beck (2016)	Compared four MIP formulations and found modern MIP effective on moderate instances, while also documenting formulation-dependent scalability.
Cebi, Atac, and Sahingoz (2020)	Reviewed exact, heuristic, and metaheuristic JSSP methods, including GA and mathematical programming.

GA and MIP are selected because they provide complementary evidence. GA follows the standard evolutionary workflow—population, chromosome, fitness, selection, crossover, mutation, and repeated generations—but can converge prematurely and cannot prove optimality. MIP expresses the objective and hard constraints directly and can prove optimality when the solver terminates with an optimal status. Comparing both on identical FT06 data separates solution quality from proof of solution quality.

2. Data Preparation and Problem Analysis

2.1 Dataset and representations

FT06 was obtained from the OR-Library [job-shop scheduling collection](#), specifically the [jobshop1 data file](#) (Beasley, 1990). The source file attributes FT06 to Fisher and Thompson (1963) and stores machine identifiers from 0 to 5. For readability, this report labels them M1 to M6. FT06 has six jobs, six machines, 36 operations, and a published optimal makespan of 55 (Weckman, Ganduri, and Koonce, 2008). The Python implementation stores the processing times and routes in the two representations required by both methods.

Processing-time matrix P :

Job	O1	O2	O3	O4	O5	O6
J1	1	3	6	7	3	6
J2	8	5	10	10	10	4
J3	5	4	8	9	1	7
J4	5	5	5	3	8	9
J5	9	3	5	4	3	1
J6	3	3	9	10	4	1

Machine-route matrix M :

Job	O1	O2	O3	O4	O5	O6
J1	M3	M1	M2	M4	M6	M5
J2	M2	M3	M5	M6	M1	M4
J3	M3	M4	M6	M1	M2	M5
J4	M2	M1	M3	M4	M5	M6
J5	M3	M2	M5	M6	M1	M4
J6	M2	M4	M6	M1	M5	M3

The full data definition and implementation are provided in [question3_generate_assets.py](#). Preprocessing converts the report's one-based labels, such as M1, to zero-based Python list indices and creates one record per operation. No imputation, scaling, or statistical transformation is appropriate because the benchmark parameters are complete deterministic integers. The same matrices supply GA decoding and MIP coefficients, ensuring a like-for-like comparison.

Assumptions are deterministic processing times, non-preemptive operations, continuous machine availability from time zero, and one operation per machine at a time. Setup times, due dates, staff limits, maintenance, and machine breakdowns are outside the selected classical JSSP.

2.2 Lower bound and complexity

The longest job requires 47 time units and the largest machine workload is 43, giving the simple lower bound

$$C_{\max} \geq \max(47, 43) = 47.$$

An operation-based chromosome has 36 positions containing six copies of each job. Its number of distinct sequences is

$$\frac{36!}{(6!)^6} \approx 2.67 \times 10^{24}.$$

Many sequences decode to the same schedule, but this figure still shows why exhaustive enumeration is unsuitable. The known optimum of 55 is used only for reporting and benchmark comparison; it never guides parent selection, crossover, mutation, decoding, hyperparameter selection, or stopping.

3. Genetic Algorithm Design, Tuning, and Results

3.1 Complete GA design

The chromosome is a length-36 list in which each job identifier occurs six times. The k -th appearance of a job schedules that job's k -th operation. Reading genes from left to right, the decoder starts each operation at

$$S_a = \max\{\text{job-ready time, machine-ready time}\}.$$

Consequently, job precedence and machine capacity are satisfied by construction. This is a domain-specific refinement of generic scheduling algorithms: standard fitness functions often penalize infeasible allocations, whereas this decoder produces only feasible schedules and can use makespan directly.

GA component	Implemented design
Population initialization	Random shuffles of six copies of each job; a local seeded random generator makes every run reproducible.
Fitness	$f(z) = 1/C_{\max}(z)$; implementation ranks the equivalent makespan directly, with lower values preferred.
Selection	Tournament selection, tournament size 3.
Crossover	Job Order Crossover (JOX), which preserves six occurrences of every job.
Mutation	Swap mutation, which exchanges two genes and preserves job counts while maintaining diversity.
Elitism	Best four chromosomes copied unchanged to the next generation.
Stopping criterion	The configured maximum number of generations; the benchmark optimum does not stop the search.
Best-solution tracking	Best-so-far chromosome, makespan, fitness, schedule, and first generation matching 55.

This sequence follows a standard GA workflow. JOX and swap mutation replace generic one-point and random-reset operators because those operators could invalidate the repeated-job chromosome.

3.2 Hyperparameter tuning

The hyperparameter tuning process distinguishes a search space, validation metric, reproducible search, and held-out evaluation. A full $3^4 = 81$ grid would be unnecessarily expensive, so 12 configurations were sampled reproducibly from that grid using search seed 2025. All values of every required hyperparameter were covered:

Hyperparameter	Values searched
Population size	50, 100, 150
Crossover probability	0.80, 0.90, 0.95
Mutation probability	0.10, 0.20, 0.30
Maximum generations	100, 300, 500

Each candidate was run on tuning seeds 11, 29, and 47. Selection minimized mean makespan, then worst makespan; remaining ties used makespan standard deviation and the nominal population-by-generation budget. The known optimum and hit rate were descriptive evaluation measures, not selection targets. Runtime was recorded but not used as the first criterion because wall-clock time varies with the execution environment. The complete 36-run tuning stage took approximately 54.4 seconds and was a one-off model-selection cost.

Here, standard deviation is the population standard deviation of the three best makespans obtained with seeds 11, 29 and 47. It is calculated with `statistics.pstdev`. It measures variation across these three tuning runs and is not a standard error or confidence interval.

ID	Pop.	p_c	p_m	Gen.	Mean	Worst	SD	Hit rate	Selected
1	150	0.90	0.30	500	55.67	57	0.943	66.7%	No
2	50	0.90	0.10	300	57.33	59	1.700	33.3%	No
3	150	0.80	0.30	300	57.67	59	1.886	33.3%	No
4	50	0.95	0.20	300	57.67	58	0.471	0.0%	No
5	150	0.90	0.20	300	55.00	55	0.000	100.0%	No
6	50	0.80	0.10	100	58.33	59	0.471	0.0%	No
7	100	0.95	0.10	500	57.00	59	1.633	33.3%	No
8	100	0.95	0.20	100	55.00	55	0.000	100.0%	Yes
9	150	0.95	0.10	100	57.33	59	1.700	33.3%	No
10	100	0.80	0.10	500	56.00	58	1.414	66.7%	No
11	50	0.80	0.30	500	58.67	59	0.471	0.0%	No
12	150	0.95	0.10	300	57.33	59	1.700	33.3%	No

Candidates 5 and 8 tied on all solution-quality criteria. Candidate 8 was selected because its nominal evaluation budget, 100×100 , was lower than candidate 5's 150×300 . Elitism and tournament size were fixed design choices, not tuned parameters.


3.3 Independent evaluation, convergence, and best solution

The selected configuration was then evaluated on five separate seeds that were not used for selection.


Measure	Evaluation result
Best makespan	55
Mean makespan	56.00
Population standard deviation of best makespan across five runs	1.265
Worst makespan	58
Runs matching the optimum	3 of 5 (60%)
Mean runtime per run	0.508 seconds
Generations per run	100

This standard deviation describes variation across the five evaluation seeds. It should not be interpreted as a confidence interval because only five runs were used.

The best evaluated run was seed 101. Its initial best makespan was 59, it first reached 55 at generation 6, and it retained 55 through generation 100 because of elitism. The final population mean was 57.45. The plot therefore shows both rapid best-solution improvement and continued population variation.



The best-run schedule below contains all 36 operations and was independently checked for correct durations and routes, non-negative starts, job precedence, machine non-overlap, and agreement between the schedule makespan and reported objective.



The best GA run has a benchmark gap of

$$\frac{55 - 55}{55} \times 100 = 0\%.$$

This establishes that GA **found** an optimal schedule in that run; it does not mean GA proved optimality or reached 55 reliably. The 60% hit rate is the relevant evidence of stochastic variability outside the tuning seeds.

3.4 Computational efficiency and scalability

Chromosome validation uses a frequency counter and decoding scans each of the nm genes once. Therefore, validation plus decoding for one chromosome is $O(nm)$. Ignoring operator overhead, the selected run performs at most

$$100 \text{ individuals} \times 100 \text{ generations} \times 36 \text{ operations} = 360,000$$

operation placements, plus the initial population. Evaluations are independent and could be parallelized. For larger instances, however, the chromosome length, population required for diversity, and generations needed for reliable search are likely to increase. GA can retain practical time limits by controlling these budgets, but it provides no optimality bound and remains sensitive to representation, operators, and seed.

4. Mixed Integer Programming Implementation and Results

4.1 PuLP implementation

The mathematical model in Section 1.2 was implemented in PuLP 3.3.2 and solved with CBC using a 120-second limit. The code creates completion-time equalities, job-precedence inequalities, binary same-machine disjunctions, and makespan inequalities directly. PuLP recorded CBC status `Optimal` with status code 1. The returned schedule was then checked by an independent validator. This validator confirms feasibility and the reported makespan, but it does not by itself prove optimality. The original FT06 result is labelled optimal because the mapped solver status is `Optimal`, the validator passes, and the objective matches the published optimum. For a solve that reaches a time limit, the native CBC termination line must also be checked because the mapped PuLP status alone can be misleading.


Model component	Count
Start-time variables S_a	36
Completion-time variables C_a	36
Binary sequencing variables x_{ab}	90
Makespan variable	1
Total variables	163
Completion equalities	36
Job-precedence inequalities	30
Machine non-overlap inequalities	180
Makespan inequalities	36
Total constraints	282

The 90 binaries arise because each of six machines processes six operations: $6 \binom{6}{2} = 90$. The valid Big-M horizon is $H = 197$.

4.2 Objective, feasibility, computation time, and best solution

Measure	MIP result
CBC status	Optimal
Objective / makespan	55
Difference from published optimum	0.00%
Runtime	2.153 seconds
Operations validated	36
Schedule validation	Passed

Validation confirmed the equality definitions, all job precedences, non-negative times, correct routes and durations, absence of machine overlaps, and a calculated makespan of 55. PuLP recorded CBC status `Optimal`, and the feasible objective of 55 matches the independently published FT06 optimum. Therefore, the returned schedule is optimal for this benchmark. The native log for this original 120-second run was not saved. Section 4.4 reports a separate fixed-budget FT06 rerun whose archived CBC log also proves an objective and bound of 55.



The stored run summary records Python 3.13.5, PuLP 3.3.2, macOS 12.7.5 and an x86_64 processor architecture. The separate scalability logs record CBC 2.10.3. The processor model and memory size were not recorded. The reported runtimes are therefore local observations and should not be treated as hardware-independent performance results.

4.3 MIP scalability

If every one of n jobs visits each of m machines once, pairwise sequencing requires

$$m \binom{n}{2} = O(mn^2)$$

binary variables and twice as many disjunctive inequalities. The branch-and-bound search can grow much faster than this model-size expression suggests. Big-M constraints may also give weak linear relaxations. Thus the FT06 result demonstrates exactness on a small instance, not guaranteed fast performance on large job shops. The next section adds a bounded FT10 experiment rather than relying only on this theoretical argument.

4.4 Fixed-budget empirical scalability check

FT10 was selected as a larger standard instance from the same OR-Library file. It has 10 jobs, 10 machines and 100 operations, compared with 6 jobs, 6 machines and 36 operations in FT06. Its published optimum is 930 (Beck and Fox, 2000). This supplementary experiment used the selected FT06 GA configuration without retuning, the same five seeds for both instances, and a 30-second single-thread CBC limit. Using the same budgets makes the comparison reproducible, although it does not show the best performance that either method could reach after instance-specific tuning.

Instance and method	Fixed budget	Makespan result	Runtime	Feasibility and optimality evidence
FT06 GA	Population 100, 100 generations, five seeds	Best 55; mean 56.0; worst 58	1.154 s mean	All best schedules validated; three runs matched 55, but GA supplied no proof
FT06 MIP	CBC, one thread, 30 s limit	Incumbent 55; bound 55	8.746 s	Native log: Optimal solution found ; validated schedule; 0% solver gap
FT10 GA	Same GA configuration and five seeds	Best 1,036; mean 1,071.4; worst 1,121	1.874 s mean	All best schedules validated; best result was 11.4% above 930; no proof
FT10 MIP	CBC, one thread, 30 s limit	Incumbent 1,147; bound 724	30.305 s	Validated incumbent; native log: Stopped on time limit ; 58% solver-reported gap

These are fresh timings from the fixed-budget experiment. They are kept separate from the original FT06 timings because solver settings and computer load can change measured runtime.

The MIP formulation grew from 90 to 450 binary variables and from 282 to 1,190 constraints. CBC proved FT06 but reached the time limit on FT10. The FT10 incumbent was feasible, but the wide bound means it was not proved near-optimal. PuLP returned status code 1 for this time-limited run, so the native CBC termination line was used to avoid incorrectly labelling it optimal.

Under the same GA budget, mean runtime increased moderately while the best objective moved 11.4% above the published optimum. In this bounded test, GA found a better FT10 incumbent than MIP. This does not establish that GA is generally better because neither method was tuned for FT10 and only two instances were tested. The result still provides practical scalability evidence: GA keeps a controlled search budget, while proving a MIP optimum becomes harder as the model and search tree grow.

5. Comparative Analysis and Critical Reflection

Required comparison	Genetic Algorithm	Mixed Integer Programming
Solution quality and objective	Best run 55; mean 56; worst 58; 60% of evaluation runs reached 55	Objective 55 in one CBC solve
Optimality evidence	Empirical 0% benchmark gap for the best run; no proof	The validated makespan of 55 matches the published optimum. The original run's native log was not archived, but the separate fixed-budget rerun proved objective 55 with bound 55.
Computational performance	Mean 0.508 s per held-out run; one-off 36-run tuning cost 54.4 s	2.153 s for model construction and solve
Scalability	Evaluation budget is controllable and parallelizable, but reliable quality may require larger populations and more generations	Binary pairs and solver search grow rapidly; proving optimality can become difficult
Complex constraints	Decoder can be extended and soft rules can enter fitness penalties, but feasibility logic must be designed and tested	Hard rules are explicit and auditable, but each added constraint enlarges or complicates the model
Flexibility and real-world applicability	Well suited to changing objectives, soft preferences, and time-limited search	Well suited when constraints are stable and a proof or bound is important

On FT06, MIP is preferable when the decision-maker requires a proven optimum and the model remains tractable. GA is preferable when instance size or operational complexity makes exact proof too costly and a high-quality feasible schedule within a controlled time is sufficient. A selected-configuration GA run was faster here, but its one-off tuning stage was slower than the MIP solve. Neither observation establishes general superiority: FT06 is small, CBC and Python have different implementation overheads, and GA displayed seed-dependent quality.

The experiment has four important limitations. First, FT06 and FT10 are deterministic benchmark proxies rather than live factory data. Second, the primary tuning and method comparison use FT06; the fixed-budget FT10 check adds empirical scaling evidence but two instances are insufficient for a general scaling curve. Third, three tuning seeds and five

evaluation seeds provide a transparent but limited robustness sample. Fourth, the GA decoder and operators represent one design; another representation or local search could change the comparison.

Possible improvements required by the analysis are:

1. **Hybrid GA–MIP:** use GA to supply a strong feasible schedule as a MIP warm start, then let MIP improve it or certify a gap.
2. **Parallel optimization:** evaluate GA chromosomes or independent seeds concurrently; MIP could also use solver parallelism where available.
3. **Other metaheuristics and stronger models:** compare tabu search or simulated annealing with GA, and test tighter pair-specific Big-M values or valid inequalities for MIP.

In conclusion, both implementations returned a feasible schedule with makespan 55, but the evidence differs. GA found that value in three of five independent runs. MIP returned a solver-reported optimum of 55, which also matched the independently published optimum. This retains the original conclusion that MIP gives stronger assurance for this small benchmark. For larger or more changeable scheduling problems, GA or a hybrid approach may offer a more practical balance between time, flexibility and solution quality.

6. Reproducibility and Supporting Files

The complete GA, MIP, validation and figure-generation code is provided in [question3_generate_assets.py](#). The fixed-budget comparison is provided in [question3_scalability_experiment.py](#). The recorded run used Python 3.13.5, PuLP 3.3.2 and CBC 2.10.3. The pinned dependency is listed in [question3_requirements.txt](#). From the project directory, a clean environment can be created and both analyses regenerated with:

```
bash
```

```
python3 -m venv .venv-q3
.venv-q3/bin/python -m pip install -r question3_requirements.txt
.venv-q3/bin/python question3_generate_assets.py
.venv-q3/bin/python question3_scalability_experiment.py
```

The stored evidence includes:

- [GA tuning results](#), containing all 12 candidates and the three-seed summary statistics;
- [GA evaluation results](#), containing the five held-out runs;
- [best GA schedule](#), containing all 36 validated operations;
- [optimal MIP schedule](#), containing all 36 validated operations;
- [run summary](#), containing the seeds, selected configuration, aggregate results, MIP status, model size, validation results and recorded software environment;
- [scalability results](#), containing the fixed-budget FT06 and FT10 comparison;
- [scalability run summary](#), containing the full GA aggregates, model sizes, schedule checks and solver evidence; and
- [FT06 CBC log](#) and [FT10 CBC log](#), preserving the native solver termination, incumbent, bound, gap and runtime evidence.

The fixed random seeds reproduce the GA search decisions. Runtime values may change across computers. A different CBC build may also return a different schedule with the same optimal makespan because FT06 can have more than one optimal schedule.

References

- Adams, J., Balas, E., and Zawack, D. (1988). The shifting bottleneck procedure for job shop scheduling. *Management Science*, 34(3), 391–401. <https://doi.org/10.1287/mnsc.34.3.391>
- Beasley, J. E. (1990). OR-Library: Distributing test problems by electronic mail. *Journal of the Operational Research Society*, 41(11), 1069–1072. <https://doi.org/10.1057/jors.1990.166>
- Beasley, J. E. (n.d.). *Job shop scheduling data: jobshop1*. OR-Library. <https://people.brunel.ac.uk/~mastjjb/jeb/orlib/files/jobshop1.txt>
- Beck, J. C., and Fox, M. S. (2000). Dynamic problem structure analysis as a basis for constraint-directed scheduling heuristics. *Artificial Intelligence*, 117(1), 31–81. [https://doi.org/10.1016/S0004-3702\(99\)00099-5](https://doi.org/10.1016/S0004-3702(99)00099-5)
- Cebi, C., Atac, E., and Sahingoz, O. K. (2020). Job Shop Scheduling Problem and Solution Algorithms: A Review. *2020 11th International Conference on Computing, Communication and Networking Technologies*, 1–7. <https://doi.org/10.1109/ICCCNT49239.2020.9225581>
- Cheng, R., Gen, M., and Tsujimura, Y. (1996). A tutorial survey of job-shop scheduling problems using genetic algorithms—I. Representation. *Computers & Industrial Engineering*, 30(4), 983–997. [https://doi.org/10.1016/0360-8352\(96\)00047-2](https://doi.org/10.1016/0360-8352(96)00047-2)
- Fisher, H., and Thompson, G. L. (1963). Probabilistic learning combinations of local job-shop scheduling rules. In J. F. Muth and G. L. Thompson (Eds.), *Industrial Scheduling* (pp. 225–251). Prentice-Hall.
- Garey, M. R., Johnson, D. S., and Sethi, R. (1976). The complexity of flowshop and jobshop scheduling. *Mathematics of Operations Research*, 1(2), 117–129. <https://doi.org/10.1287/moor.1.2.117>
- Gonçalves, J. F., Mendes, J. J. M., and Resende, M. G. C. (2005). A hybrid genetic algorithm for the job shop scheduling problem. *European Journal of Operational Research*, 167(1), 77–95. <https://doi.org/10.1016/j.ejor.2004.03.012>
- Ku, W.-Y., and Beck, J. C. (2016). Mixed Integer Programming models for job shop scheduling: A computational analysis. *Computers & Operations Research*, 73, 165–173. <https://doi.org/10.1016/j.cor.2016.04.006>
- Weckman, G. R., Ganduri, C. V., and Koonce, D. A. (2008). A neural network job-shop scheduler. *Journal of Intelligent Manufacturing*, 19(2), 191–201. <https://doi.org/10.1007/s10845-008-0073-9>